

Interview Transcript – Participant A

Study: Early Value Challenges for Non-Technical Users of AI Automation Builders

Date: 23 March 2026

Duration: ~30 minutes

Interviewer: J (researcher)

Participant: A (pseudonym)

Participant Role: Civil servant at a national government department in Canada, working in an administrative case management capacity involving vendor coordination and language training facilitation. No formal technical or programming background — undergraduate degree in humanities. Self-taught AI and automation user pursuing independent AI research as a personal project outside of work, including a submitted conference paper. Describes himself as a non-technical user who relies entirely on AI coding assistants to conduct research and build software.

Anonymization note: All identifying information (name, specific department, university, specific tool brand names where used as identifiers) has been removed or replaced with neutral descriptors.

J: Yeah, all right, let's get started. Yeah, so could you briefly introduce yourself and what your day-to-day responsibilities are?

A: Sure, so I'm a case manager on behalf of a national government department. What that means is the most complicated thing I do on a day-to-day basis is fairly routine administrative work, but I've been moonlighting, so to speak, as an AI researcher. I am currently waiting on reviews from a major machine learning conference, so I should be getting my first peer reviews in just over 48 hours.

J: Nice, yeah, so how would you describe your technical background? Do you have any experience with programming or building things digitally?

A: Practically none. I did some light coding back in high school. My post-secondary education was in history — a Bachelor of Arts — but apart from that, just what I've dabbled in over the years. I have no formal training, if that answers the question.

J: Okay, yes, it does answer it. Yeah, so you already mentioned using some automation tools, so which tools did you use and what made you use them or try them in the first place?

A: So in particular, I started back in August. I ran a kind of red teaming game that I mostly vibe coded by prompting an AI assistant, and then it was about that time that a couple of AI coding agents converged in terms of capability. Since those came out in late August or September, I've used those tools almost exclusively. I barely write code anymore at all. They're essentially automating all of my experimental pipelines and my analysis scripts with natural language processing.

J: Okay, nice. Yeah, so what were you actually trying to accomplish when you first used the tools? What were you trying to automate, or what was the goal initially?

A: All kinds of different things. I am trying to escape my day job, so to speak, so I guess that's the motivation. But essentially, it became just a much faster, more efficient way to get things done when you can just throw a code repository at the tool and have it produce pull requests for you that you just click merge. You do have to have a look at things every so often because they can go awry, but by and large they're pretty solid and reliable. With the cross-checking you can do between tools and

with a large language model, it's fairly effective as long as you're just a little bit careful. They seem to have gotten even better since I would have started using them in August, and I seldom have any issues now.

J: Okay, and before you opened the tool for the first time, what did you think it would be able to do for you? What were your expectations?

A: I couldn't say that I even had any expectations, because I had never worked with these tools before. I had some vague awareness of version control prior to August 2025, and I think my first commits start basically in August 2025, so I had no idea. Since then, I've learned quite a lot. There's a bit of an off period in December, but since then it's pretty consistent.

J: Nice, okay. Yeah, so if you think back to your very first session, could you walk me through what happened step by step?

A: Gosh, it was so long ago. I don't even remember what I was working on. Probably it would have been a language tutor idea that never really got off the ground. Because my day job involves linking people up with vendors in order to learn second official languages — my government clients need to be able to speak both official languages, and so I arrange contracts between vendors and individuals and we pay all bills for those individuals. I realized that using a language model would be a lot cheaper and provide a lot more accessibility for people, because one-on-one language instruction is very expensive and supply is constrained. So I proposed a pilot of a language tutor concept — they didn't take me up on it, I wish they had — but I was able to scaffold up a functional web app. It was basically a wrapper that split the cognition between three model calls as opposed to just one. That's one of the mistakes a lot of LLM wrappers make — they just throw everything at one model, and that has issues when the LLM has to split its tasks. So splitting the tasks into three seemed to work fairly well. It was just a matter of prompting the model, telling it what I wanted, telling it how the app should work, and then iterating and testing. The tool would basically write all the code, and when it came time, I'd hit merge, and then I would try it and say, this isn't working, or I need to see this for testing, and slowly but surely we iterated on what was working. After that I tried a couple other applications, and then in November I started exploring whether models have a personality, and about three to four weeks later I was submitting a paper to a preprint archive. That same finding now exists as a preprint, so that's kind of what happened.

J: No, that's great, thank you. Yeah, so for understanding the system — when you started using it, did you understand what the tool was doing the whole time, or did you have moments where you weren't sure?

A: Pretty much, I understood. It's fairly transparent — the tool reads your existing code, it reads any instructions that it finds there, and then it makes changes with various behind-the-scenes scaffolding. It's not particularly difficult to understand what it's doing and how it works. It's kind of a closed loop until it's got something for you, and then it proposes a bunch of code and things that you can look at.

J: Okay, and did you also really feel in control of what was happening, or were you handing things over sometimes without knowing what would happen?

A: I did have some friction initially when I would use natural language prompting and the tool would infer what I meant, and then write a bunch of heuristics when that was not what I wanted. So there

certainly was some learning frustration and some learning pain as part of the process.

J: Yeah, so how would you describe that learning curve? Did it take a long time before you really got ahead of the curve?

A: I would say not a long time, but after several code changes and realizing that the tool had gone very much awry, I learned rather painfully. I guess it was a little bit steep, but that's because — I totally understand the failure mode — it wasn't really the tool's fault. It was my fault for not prompting the tool correctly, or not catching and reviewing things carefully enough. That's the kind of thing that I think could happen. I still wonder what to do with that particular project, but I've since moved on to do a couple other things, so I'm not too bitter about it.

J: Yeah, okay. Yeah, so when you got stuck, how did you try to get unstuck? Did you use documentation, for example, or was it trial and error? What's your go-to method?

A: Trial and error and a lot of brainstorming with an AI chat assistant — basically showing it the code and saying, okay, this happened, how do we get out of this? Using one AI agent to fix another agent's mistakes or issues.

J: Yeah, okay, and did it happen quite often that you got stuck, or does it go smoothly most of the time?

A: Ever since that major issue I had, it's been pretty smooth since. But of course I'm more careful now about telling it more precisely what I want it to do. And I think recent model changes have probably made them just a little bit more cautious as well.

J: Yeah, okay, yeah. And were there any other challenges that you walked into when you used AI automation tools?

A: Nothing beyond your standard iteration — something doesn't work and you fix it. I think that one of the issues with AI in general is that it can be very credulous. When I discovered my research finding, one of the models I was using was really, really credulous in terms of iteration and discussion — it just believed everything. And I kind of went, no, that doesn't make a lot of sense. The results didn't make sense, and we eventually had to apply a logarithmic transformation in the math to make it work. So that was an issue, but it was not a major issue — it was definitely on me to catch that the results just didn't make a lot of sense.

J: Okay, yeah, clear. Yeah, so you don't really remember your first session, but do you remember feeling like you had accomplished something useful in maybe the first few sessions?

A: Oh, yeah, absolutely, absolutely. AI is allowing me to do things that I could never do on my own. I'm not a coder by training or anything like it, and yet I'm sitting on a conference review system right now waiting for my reviews to come in. That's for someone who has no computer science background whatsoever, no mathematics background, no true technical background — apart from being a certain generation that grew up with the internet partially in their lives. I'm doing things today that would be impossible without AI automation — applying stats, using models. It's truly exhilarating to think where we could go. Automation is one thing, but I prefer augmentation. It's an augment for my cognitive abilities.

J: Yeah, that's a nice way to look at it. Yeah, so if the tools could have done one thing differently to make the first experience better, do you have any idea of what that could be?

A: Yeah, I do. One of the tools I use defaults to asking me questions to clarify what it should do, and the other is very deterministic and just makes assumptions. Clarity beats assumptions every time. A lot of my issues happened from assumptions and me being imprecise. So if it had asked me to clarify what I wanted, we wouldn't have run into the mess that we did.

J: Yeah, okay. Yeah, so did this whole experience change how willing you are to try automating other things in your work?

A: Oh, absolutely, absolutely. I wouldn't even say that I'm necessarily slow to adopt. I'm not even really slow to learn. I haven't tried much else right now because for what I'm doing, this seems sufficient for the time being, until something new comes along. I've downloaded and played a little bit with a local model runner, but I haven't used it extensively — just a little bit to try and get used to the architecture and see what kinds of things it could do for me. Of course, I burned out my GPU, and I wanted to use a local model for inference, so I'm now in the market for a new GPU. In local currency it's quite a significant investment and you don't always know where the money's going to come from. But it'll get done eventually.

J: Yeah, okay. Yeah, so when you tried using that local model runner, what did you try to do? I'm really curious.

A: Automate my experiments. At the time I was working on a large benchmark — over 400 scenarios for an AI safety-related research project, and each scenario was subjected to chain-of-thought decompositions, so there were nearly 900 prompts. When inference takes like a week to run through, you don't want to babysit the computer for weeks. So it would have been nice to have the local agent automatically start the next experiment once the current one finished. That didn't end up happening for different reasons, but that was what I was hoping to do with it.

J: Nice, and was it easy to get started with the local model runner, or was it difficult to set up?

A: No, because in my experience, it's just not clear how it's meant to work. They really funnel you into a singular process of doing it the way that they think it should be done, and so no, regrettably, it was not as easy. Of course, there have probably been hundreds of updates since, and I haven't played with it for probably six or eight weeks now, so it's been a bit of time.

J: Yeah, okay, cool. Because I've also tried to do some things with it, but I thought it was a bit of a time investment to really get it working. I wanted to automate something of the development as well for our startup, but it was a bit time consuming to set up — I wanted to put it on a server and just let it develop basically like a junior developer, and then our developer in our team could look at it as a senior developer. But yeah, it was a bit difficult to set up. Very interesting though. Yeah, I think I already got some really helpful and useful insights, so thank you for that. Do you have any other things to add, or anything that crossed your mind that might be interesting?

A: I think what strikes me is the cognitive unlock that AI is capable of doing for people is being left on the table by the overwhelming majority of users. A lot of people are still learning prompt engineering, and while that's kind of important, the natural language conversations you can have are much more capable these days than they used to be. The long context conversations can really help people, and I hope that recognition picks up, because what we can do is so much greater than what a lot of people think is possible now. I'm also concerned about how much fear the general public has of these tools. There's a paradigm where people think AI will take over the world. It might

— I'm not saying it's all wrong — but the issue is that it's not clear to me that simply scaling models will suddenly produce volition. My benchmarks don't suggest that. That's not to say there aren't problems, but I think the problems are far more in keeping with humans doing bad things to each other than AI doing its own thing.

J: Yeah, interesting. Yeah, I'm kind of out of my questions. I think I got a lot of insights already. Do you have any questions for me?

A: No, that's all. I don't have any questions for you. Thank you.

J: All right. Thank you very much for your time. And if you might know someone else who's also a non-technical user and has tried AI automation tools, I would be very happy to talk to them. So yeah, just let me know. And yeah, if you also want to be notified of the progress of my thesis, then I can also keep you in the loop.

A: I'd love to be kept in the loop.

J: I think that'd be great. Yeah, I will keep you posted. All right. Thank you very much. I really got some useful insights. So thank you very much for your time.

A: You too. Bye-bye now.

J: Bye-bye.